

B-TREES

Alejandro Martín
Juan Mendivelso

Contents

1. Introduction
2. Definition of B-Trees
3. Basic operations on B-Trees
4. Deleting a key from a B-Tree

1. INTRODUCTION

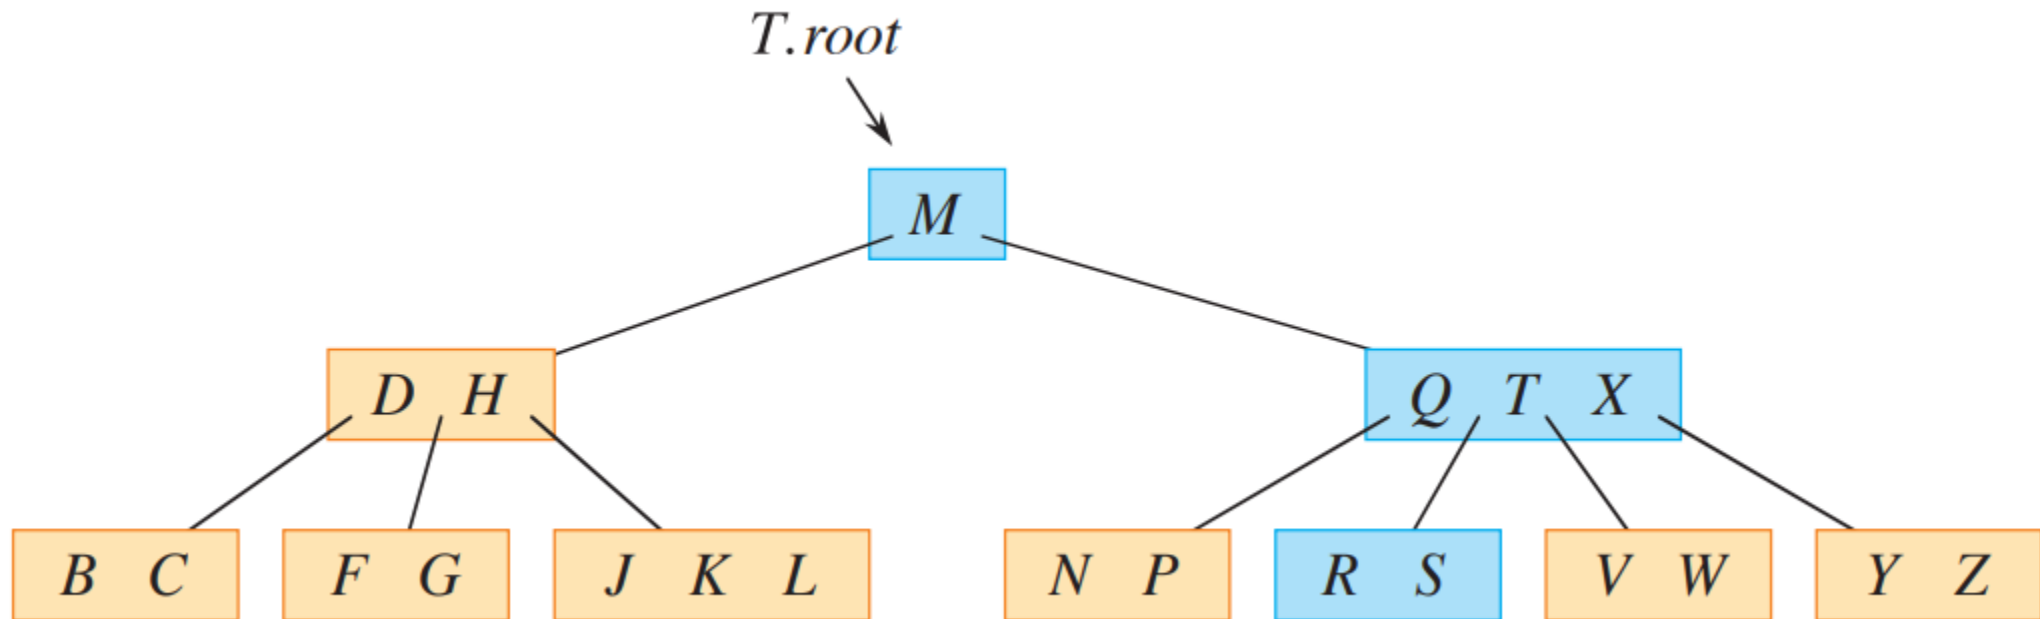
B-Trees

- B-Trees are balanced search trees designed to work well on disk drives or other direct-access secondary storage devices.
- Similar to Red-Black trees, but they are better at minimizing the number of operations that access disks.
- B-Tree nodes may have many children from few to thousands, contrary to Red-Black trees.
- Every B-Tree has height $O(\lg n)$. But a B-Tree has a larger branching factor than a Red-Black tree.

B-Trees


- B-Trees generalize binary search trees in a natural manner.
- If an internal B-tree node x contains $x.n$ keys, then x has $x.n + 1$ children.
- The keys in node x serve as dividing points separating the range of keys handled by x into $x.n + 1$ subranges, each handled by one child of x .
- A search for a key in a B-tree makes an $(x.n + 1)$ -way decision based on comparisons with the $x.n$ keys stored at node x .

B-Trees



Data structures on secondary storage

- Computer systems take advantage of various technologies that provide memory capacity:
 - Main memory
 - Secondary storage (Magnetic disk drives, SSD)



Data structures on secondary storage

- In order to amortize the time spent waiting for mechanical movements (latency), disk drives access more than one item at a time.
- Information is divided into a number of equal-sized blocks of bits that appear consecutively within tracks, and each disk read or write is one or more entire blocks.

Data structures on secondary storage

- Often, accessing a block of information and reading it from a disk drive takes longer than processing all the information read.
- We will look separately at the two principal components of the running time:
 - The number of disk accesses
 - The CPU (computing) time.

Data structures on secondary storage

- We measure the number of disk accesses in terms of the number of blocks of information that need to be read from or written to the disk drive.
- In a typical B-tree application, the amount of data handled is so large that all the data do not fit into main memory at once.
- The B-tree algorithms copy selected blocks from disk into main memory as needed and write back onto disk the blocks that have changed.
- B-Tree algorithms keep only a constant number of blocks in main memory at any time, and thus the size of main memory does not limit the size of B-trees that can be handled.

Data structures on secondary storage

- B-tree procedures need to be able to read information from disk into main memory and write information from main memory to disk.
- Consider some object x . If x is currently in the computer's main memory, then the code can refer to the attributes of x as usual.
- If x resides on disk, however, then the procedure must perform the operation *DISK-READ*(x) to read the block containing x into main memory before it can refer to x 's attributes.
- Similarly, procedures call *DISK-WRITE*(x) to save any changes that have been made to the attributes of object x by writing to disk the block containing x .

Data structures on secondary storage

x = a pointer to some object

DISK-READ(x)

operations that access and/or modify the attributes of x

DISK-WRITE(x) // omitted if no attributes of x were changed

other operations that access but do not modify attributes of x

2. DEFINITION OF B-TREES

Definition of B-Trees

- A **B-tree** T is a rooted tree with root $T.root$ having the following properties:
 1. Every node x has the following attributes:
 - a) $x.n$, the number of keys currently stored in node x ,
 - b) The $x.n$ keys themselves, stored in monotonically increased order.
 - c) $x.leaf$, a Boolean value that is TRUE if x is a leaf and FALSE if x is an internal node.
 2. Each internal node x also contains $x.n + 1$ pointers $x.c_1, x.c_2, \dots, x.c_{x.n+1}$ to its children. Leaf nodes have their c_i attributes undefined.
 3. The keys $x.key_i$ separate the ranges of keys stored in each subtree: if k_i is any key stored in the subtree with root $x.c_i$, then

$$k_1 \leq x.key_1 \leq k_2 \leq x.key_2 \leq \dots \leq x.key_{x.n} \leq k_{x.n+1}$$

Definition of B-Trees

- A **B-tree** T is a rooted tree with root $T.root$ having the following properties:
 4. All leaves have the same depth, which is the tree's height h .
 5. Nodes have lower and upper bounds on the number of keys they can contain, expressed in terms of a fixed integer $t \geq 2$ called the **minimum degree** of the B-tree:
 - a) Every node other than the root must have at least $t - 1$ keys. Every internal node other than the root thus has at least t children. If the tree is nonempty the root must have at least one key.
 - b) Every node may contain at most $2t - 1$ keys. Therefore, an internal node may have at most $2t$ children. We say that a node is **full** if it contains exactly $2t - 1$ keys.
- The simplest B-tree occurs when $t = 2$. Every internal node then has either 2, 3 or 4 children, and it is a **2-3-4 tree**.

The height of a B-Tree

- The number of disk accesses required for most operations on a B-tree is proportional to the height of the B-tree. The following theorem bounds the worst-case height of a B-Tree.
- **Theorem:** if $n \geq 1$, then for any n -key B-tree T of height h and minimum degree $t \geq 2$,

$$h \leq \log_t \frac{n+1}{2}$$

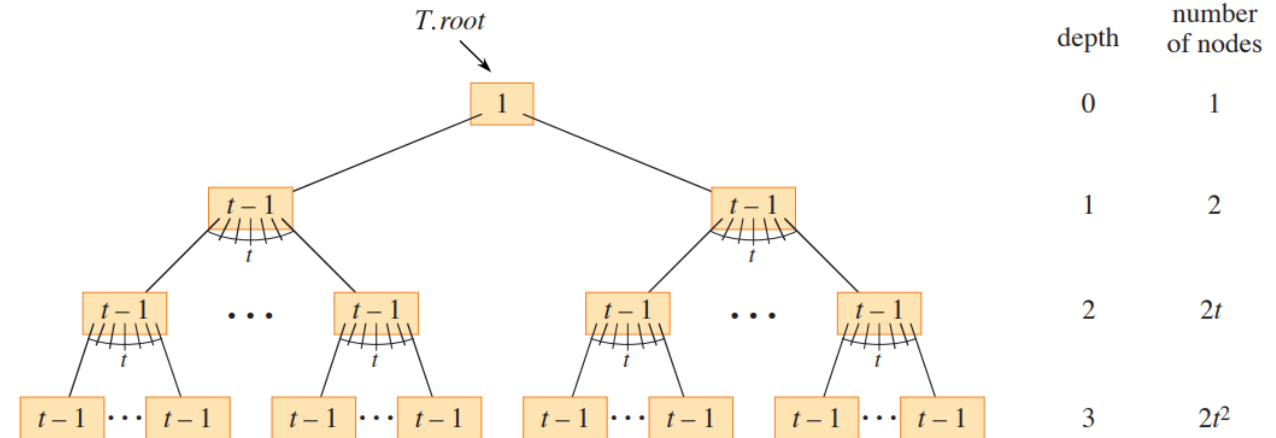
- **Proof:**
 - The root of a nonempty B-tree T contains at least one key, and all other nodes contain at least $t - 1$ keys.
 - Then T contains at least 2 nodes at depth 1, at least $2t$ nodes at depth 2, at least $2t^2$ nodes at depth 3 and so on, until at depth h it has at least $2t^{h-1}$ nodes.

The height of a B-Tree

- **Proof:**

- The number n of keys therefore satisfies the inequality:

$$\begin{aligned} n &\geq 1 + (t - 1) \sum_{i=1}^h 2t^{i-1} \\ &= 1 + 2(t - 1) \left(\frac{t^h - 1}{t - 1} \right) \\ &= 2t^h - 1 \end{aligned}$$



- So that $t^h \leq (n + 1)/2$. Then we take base- t logarithm of both sides and we finish

3. BASIC OPERATION ON B-TREES

Basic operations on B-Trees

- We present the operations B-TREE-SEARCH, B-TREE-CREATE and B-TREE-INSERT. These procedures observe two conventions:
 - The root of the B-Tree is always in main memory, so that no procedure ever needs to perform a DISK-READ on the root. However, if any changes occur in the root node, then DISK-WRITE must be called there.
 - Any nodes that are passed as parameters must already have had a DISK-READ operation performed on them.

Searching a B-Tree

- Similar like searching a BST, except that at each internal node x , the search makes an $(x.n + 1)$ -way branching decision (instead of the two-way).
- The procedure B-TREE-SEARCH generalizes the procedure defined for BST.

```
B-TREE-SEARCH( $x, k$ )
1   $i = 1$ 
2  while  $i \leq x.n$  and  $k > x.key_i$ 
3       $i = i + 1$ 
4  if  $i \leq x.n$  and  $k == x.key_i$ 
5      return  $(x, i)$ 
6  elseif  $x.leaf$ 
7      return NIL
8  else DISK-READ( $x.c_i$ )
9      return B-TREE-SEARCH( $x.c_i, k$ )
```

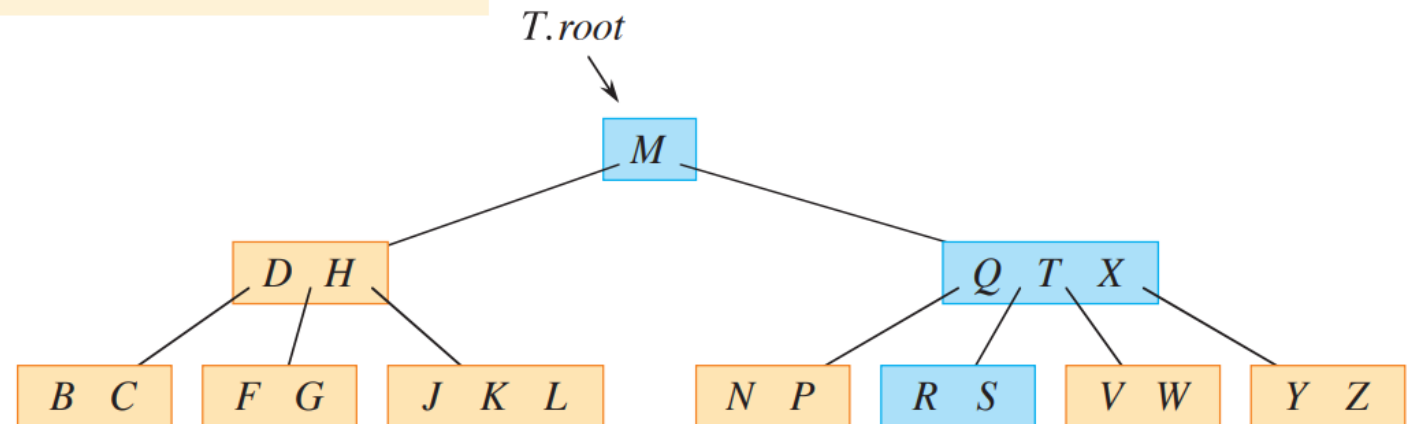
$O(h)$ disk
accesses and
 $O(th)$ CPU time

Searching a B-Tree

B-TREE-SEARCH(x, k)

```
1   $i = 1$ 
2  while  $i \leq x.n$  and  $k > x.key_i$ 
3       $i = i + 1$ 
4  if  $i \leq x.n$  and  $k == x.key_i$ 
5      return  $(x, i)$ 
6  elseif  $x.leaf$ 
7      return NIL
8  else DISK-READ( $x.c_i$ )
9      return B-TREE-SEARCH( $x.c_i, k$ )
```

B-TREE-SEARCH(R,T)



Creating an empty B-Tree

- We use the B-TREE-CREATE procedure to create an empty root node and then call the B-TREE-INSERT procedure to add new keys
- Both procedures use an auxiliary procedure (ALLOCATE-NODE) that allocates one disk block to be used as a new one in $O(1)$ time.

B-TREE-CREATE(T)

```
1   $x = \text{ALLOCATE-NODE}()$ 
2   $x.\text{leaf} = \text{TRUE}$ 
3   $x.n = 0$ 
4   $\text{DISK-WRITE}(x)$ 
5   $T.\text{root} = x$ 
```

It takes $O(1)$
time

Inserting a key into a B-Tree

- With a B-Tree, you cannot simply create a new leaf node and insert it. Instead, we insert the new key into an existing leaf node.
- Since we cannot insert a key into a leaf node that is full, we need an operation that splits a full node y around its **median key** into two nodes having only $t - 1$ keys each.
- The median key moves up into y 's parent to identify the dividing point between the two new trees.
- But if y 's parent is also full, then we must split it before we can insert the new key.
- To avoid having to go back up the tree, we just split every full node we encounter as we go down the tree

Splitting a node in a B-Tree

- The procedure B-TREE-SPLIT-CHILD takes as input a nonfull internal node x and an index i such that $x.c_i$ is a full child of x

```
B-TREE-SPLIT-CHILD( $x, i$ )  
1   $y = x.c_i$                                 // full node to split  
2   $z = \text{ALLOCATE-NODE}()$                       //  $z$  will take half of  $y$   
3   $z.\text{leaf} = y.\text{leaf}$   
4   $z.n = t - 1$   
5  for  $j = 1$  to  $t - 1$                           //  $z$  gets  $y$ 's greatest keys ...  
6       $z.\text{key}_j = y.\text{key}_{j+t}$   
7  if not  $y.\text{leaf}$   
8      for  $j = 1$  to  $t$                             // ... and its corresponding children  
9           $z.c_j = y.c_{j+t}$   
10  $y.n = t - 1$                                 //  $y$  keeps  $t - 1$  keys  
11 for  $j = x.n + 1$  downto  $i + 1$                 // shift  $x$ 's children to the right ...  
12      $x.c_{j+1} = x.c_j$   
13  $x.c_{i+1} = z$                                 // ... to make room for  $z$  as a child  
14 for  $j = x.n$  downto  $i$                         // shift the corresponding keys in  $x$   
15      $x.\text{key}_{j+1} = x.\text{key}_j$   
16  $x.\text{key}_i = y.\text{key}_t$                             // insert  $y$ 's median key  
17  $x.n = x.n + 1$                                 //  $x$  has gained a child  
18 DISK-WRITE( $y$ )  
19 DISK-WRITE( $z$ )  
20 DISK-WRITE( $x$ )
```


It takes $\Theta(t)$ CPU time and $O(1)$ disk operations

Splitting a node in a B-Tree

B-TREE-SPLIT-CHILD(x, i)

```

1   $y = x.c_i$                                 // full node to split
2   $z = \text{ALLOCATE-NODE}()$                     //  $z$  will take half of  $y$ 
3   $z.\text{leaf} = y.\text{leaf}$ 
4   $z.n = t - 1$ 
5  for  $j = 1$  to  $t - 1$                         //  $z$  gets  $y$ 's greatest keys ...
6       $z.\text{key}_j = y.\text{key}_{j+t}$ 
7  if not  $y.\text{leaf}$ 
8      for  $j = 1$  to  $t$                         // ... and its corresponding children
9           $z.c_j = y.c_{j+t}$ 
10  $y.n = t - 1$                                 //  $y$  keeps  $t - 1$  keys
11 for  $j = x.n + 1$  downto  $i + 1$             // shift  $x$ 's children to the right ...
12      $x.c_{j+1} = x.c_j$ 
13  $x.c_{i+1} = z$                                 // ... to make room for  $z$  as a child
14 for  $j = x.n$  downto  $i$                     // shift the corresponding keys in  $x$ 
15      $x.\text{key}_{j+1} = x.\text{key}_j$ 
16  $x.\text{key}_i = y.\text{key}_t$                         // insert  $y$ 's median key
17  $x.n = x.n + 1$                                 //  $x$  has gained a child
18 DISK-WRITE( $y$ )
19 DISK-WRITE( $z$ )
20 DISK-WRITE( $x$ )
    
```



Splitting a node with $t = 4$

Inserting a key into a B-Tree in a single pass down the tree

- We use procedure B-TREE-INSERT to insert a key k into a B-Tree T of height h

```
B-TREE-INSERT( $T, k$ )  
1   $r = T.root$   
2  if  $r.n == 2t - 1$   
3       $s = \text{B-TREE-SPLIT-ROOT}(T)$   
4      B-TREE-INSERT-NONFULL( $s, k$ )  
5  else B-TREE-INSERT-NONFULL( $r, k$ )
```


It takes $\Theta(th)$ CPU time and $O(h)$ disk accesses

Inserting a key into a B-Tree in a single pass down the tree

- If the root is full, we use the procedure B-TREE-SPLIT-ROOT to split it (this is the only way to increase the height of the B-Tree).

B-TREE-SPLIT-ROOT(T)

```
1  $s = \text{ALLOCATE-NODE}()$ 
2  $s.\text{leaf} = \text{FALSE}$ 
3  $s.n = 0$ 
4  $s.c_1 = T.\text{root}$ 
5  $T.\text{root} = s$ 
6 B-TREE-SPLIT-CHILD( $s, 1$ )
7 return  $s$ 
```



Splitting the root with $t = 4$

Inserting a key into a B-Tree in a single pass down the tree

- The auxiliary procedure B-TREE-INSERT-NONFULL inserts key k into node x , which is assumed to be nonfull when the procedure is called.
- It recurses as necessary down the tree, at all time guaranteeing that the node to which it recurses is not full by calling B-TREE-SPLIT-CHILD if necessary.

B-TREE-INSERT-NONFULL(x, k)

```
1   $i = x.n$ 
2  if  $x.leaf$                                 // inserting into a leaf?
3      while  $i \geq 1$  and  $k < x.key_i$           // shift keys in  $x$  to make room for  $k$ 
4           $x.key_{i+1} = x.key_i$ 
5           $i = i - 1$ 
6       $x.key_{i+1} = k$                           // insert key  $k$  in  $x$ 
7       $x.n = x.n + 1$                           // now  $x$  has 1 more key
8      DISK-WRITE( $x$ )
9  else while  $i \geq 1$  and  $k < x.key_i$         // find the child where  $k$  belongs
```

```
10       $i = i - 1$ 
11       $i = i + 1$ 
12      DISK-READ( $x.c_i$ )
13      if  $x.c_i.n == 2t - 1$                   // split the child if it's full
14          B-TREE-SPLIT-CHILD( $x, i$ )
15          if  $k > x.key_i$                     // does  $k$  go into  $x.c_i$  or  $x.c_{i+1}$ ?
16               $i = i + 1$ 
17      B-TREE-INSERT-NONFULL( $x.c_i, k$ )
```

Inserting a key into a B-Tree in a single pass down the tree

B-TREE-INSERT(T, k)

```


1   $r = T.root$ 
2  if  $r.n == 2t - 1$ 
3       $s = \text{B-TREE-SPLIT-ROOT}(T)$ 
4      B-TREE-INSERT-NONFULL( $s, k$ )
5  else B-TREE-INSERT-NONFULL( $r, k$ )
    
```

B-TREE-INSERT-NONFULL(x, k)


```

1   $i = x.n$ 
2  if  $x.leaf$                                 // inserting into a leaf?
3      while  $i \geq 1$  and  $k < x.key_i$         // shift keys in  $x$  to make room for  $k$ 
4           $x.key_{i+1} = x.key_i$ 
5           $i = i - 1$ 
6       $x.key_{i+1} = k$                         // insert key  $k$  in  $x$ 
7       $x.n = x.n + 1$                         // now  $x$  has 1 more key
8      DISK-WRITE( $x$ )
9  else while  $i \geq 1$  and  $k < x.key_i$         // find the child where  $k$  belongs
10      $i = i - 1$ 
11      $i = i + 1$ 
12     DISK-READ( $x.c_i$ )
13     if  $x.c_i.n == 2t - 1$                   // split the child if it's full
14         B-TREE-SPLIT-CHILD( $x, i$ )
15         if  $k > x.key_i$                     // does  $k$  go into  $x.c_i$  or  $x.c_{i+1}$ ?
16              $i = i + 1$ 
17     B-TREE-INSERT-NONFULL( $x.c_i, k$ )
    
```


(a) initial tree



(b) B inserted



(c) Q inserted



The minimum degree for this B-Tree is $t = 3$

Inserting a key into a B-Tree in a single pass down the tree

B-TREE-INSERT(T, k)

```


1   $r = T.root$ 
2  if  $r.n == 2t - 1$ 
3       $s = \text{B-TREE-SPLIT-ROOT}(T)$ 
4      B-TREE-INSERT-NONFULL( $s, k$ )
5  else B-TREE-INSERT-NONFULL( $r, k$ )
    
```

B-TREE-INSERT-NONFULL(x, k)


```

1   $i = x.n$ 
2  if  $x.leaf$                                 // inserting into a leaf?
3      while  $i \geq 1$  and  $k < x.key_i$           // shift keys in  $x$  to make room for  $k$ 
4           $x.key_{i+1} = x.key_i$ 
5           $i = i - 1$ 
6       $x.key_{i+1} = k$                           // insert key  $k$  in  $x$ 
7       $x.n = x.n + 1$                           // now  $x$  has 1 more key
8      DISK-WRITE( $x$ )
9  else while  $i \geq 1$  and  $k < x.key_i$           // find the child where  $k$  belongs
10      $i = i - 1$ 
11      $i = i + 1$ 
12     DISK-READ( $x.c_i$ )
13     if  $x.c_i.n == 2t - 1$                     // split the child if it's full
14         B-TREE-SPLIT-CHILD( $x, i$ )
15         if  $k > x.key_i$                         // does  $k$  go into  $x.c_i$  or  $x.c_{i+1}$ ?
16              $i = i + 1$ 
17     B-TREE-INSERT-NONFULL( $x.c_i, k$ )
    
```


(c) Q inserted



(d) L inserted



(e) F inserted



The minimum degree for this B-Tree is $t = 3$




4. DELETING A KEY FROM A B-TREE

Deleting a key from a B-Tree

- It is analogous to insertion but more complicated. Because we can delete a key from any node, so that we must rearrange the node's children if needed.
- Just as a node should not get too big due to insertion, a node must not get too small during deletion.
- The procedure B-TREE-DELETE deletes the key k from the subtree rooted at x .
- B-TREE-DELETE prevents any node from becoming underfull (i.e. having fewer than $t - 1$ keys) while also making a single pass down the tree, searching for and deleting the key.

Deleting a key from a B-Tree

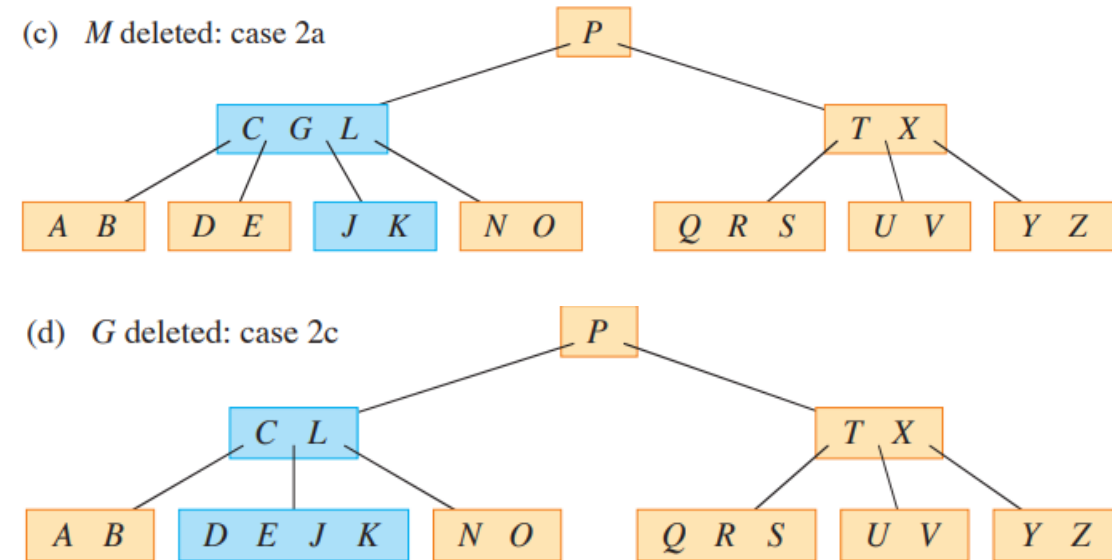
- We explain how the procedure works:
 - **Case 1:** The search arrives at a leaf node. If x contains key k , then delete k from x . If x does not contain key k , then k was not in the B-Tree, then do nothing.
 - **Case 2:** The search arrives at an internal node x that contains key k . Let $k = x.key_i$, $x.c_i$ the child that precedes k and $x.c_{i+1}$ the child that follows k . Then 3 subcases arise:
 - **Case 2a:** $x.c_i$ has at least t keys. Find the predecessor k' of k in the subtree rooted at $x.c_i$. Recursively delete k' from $x.c_i$, and replace k by k' in x .



The minimum degree for this B-Tree is $t = 3$

Deleting a key from a B-Tree





- We explain how the procedure works:
 - **Case 2:** x is internal node and contains k . Let $k = x.key_i$, $x.c_i$ child that precedes k and $x.c_{i+1}$ child that follows k :
 - **Case 2b:** $x.c_i$ has $t-1$ keys and $x.c_{i+1}$ has at least t keys. Find the successor k' of k in the subtree rooted at $x.c_{i+1}$. Recursively delete k' from $x.c_{i+1}$ and replace k by k' in x .
 - **Case 2c:** Both $x.c_i$ and $x.c_{i+1}$ have $t-1$ keys. Merge k and all $x.c_{i+1}$ into $x.c_i$, so that x loses both k and the pointer to $x.c_{i+1}$, and $x.c_i$ now contains $2t-1$ keys. Then free $x.c_{i+1}$ and recursively delete k from $x.c_i$.



The minimum degree for this B-Tree is $t = 3$

Deleting a key from a B-Tree

- We explain how the procedure works:
 - Case 3:** The search arrives at an internal node x that does not contain key k . Ensure that each node visited has at least t keys. To do so, determine the root $x.c_i$ of the appropriate subtree that must contain k . If $x.c_i$ only has $t - 1$ keys 2 cases arise:
 - Case 3a:** $x.c_i$ has only $t - 1$ keys but has an immediate sibling with at least t keys. Give $x.c_i$ an extra key by moving a key from x into $x.c_i$, moving a key from $x.c_i$'s immediate left or right sibling up into x , and moving the appropriate child pointer from the sibling into $x.c_i$.
 - Case 3b:** $x.c_i$ and each of $x.c_i$'s immediate siblings have $t - 1$ keys. Merge $x.c_i$ with one sibling which involves moving a key from x down into the new merged node to become the median key for that node.



The minimum degree for this B-Tree is $t = 3$

Deleting a key from a B-Tree

- In cases 2c and 3b, if node x is the root, it could end up having no keys. When this occurs, x is deleted and x 's only child $x.c_i$ becomes the new root of the tree.
- This procedure involves only $O(h)$ disk operations for a B-Tree of height h . The CPU time required is $O(th)$.