

Search Trees

Juan Mendivelso

Search Tree

- Structure that stores objects, each object identified by a key value, in a tree structure.
- The keys are drawn from an ordered set.
- Two keys can be compared in constant time.
- These comparisons are used to guide the access to a specific object by its key.
- The search starts from the root. Each node contains a key that is compared with the query key.
- One can go to different nodes depending on whether the query key is smaller or larger than the key in the node.

Search Tree

- This structure is fundamental to most data structure.
- It allows many variations.
- It is a building block for most more complex data structures.
- Search trees are one method to implement dictionaries:.
- A dictionary is a structure that stores objects, identified by keys, and supports the operations find, insert and delete.
- There are other ways to implement dictionaries, like hash tables.
- Also, there are applications of search trees different from dictionaries.

CONTENTS

1. Two Models of Search Trees
2. Leaf Trees
3. Operations on Leaf Trees
4. Returning from Leaf to Root
5. Dealing with Nonunique Keys
6. Queries for the Keys in an Interval
7. Building Optimal Search Trees
8. Converting a Tree into a Sorted Lists
9. Removing a Tree

1. TWO MODELS OF SEARCH TREES

Two Models of Search Tree

1. Leaf Tree

- Take left branch if query key is smaller than node key; otherwise take the right branch.
- Repeat until a leaf is found.
- The keys in the interior nodes are only for comparison.
- All objects are in the leaves.

2. Node Tree

- Take left branch if query key is smaller than node key.
- Take the right branch if the query key is larger than the node key.
- Take the object contained in the node if they are equal.

Two Models of Search Tree

1. Leaf Tree

- Each interior node has a left and a right subtree.
- It is a full binary tree: the degree of each node is either 0 or 2.
- The structure is more regular.
- One comparison for two possible outcomes: left or right.

2. Node Tree

- Left or right subtree may be missing.
- Two comparisons for three possible outcomes: node, left or right.

Two Models of Search Tree

1. Leaf Tree

- Interior nodes serve only for comparisons and may reappear in the leaves for the identification of the objects.
- It is possible that there are keys used for comparison that do not belong to any object, for example if the object has been deleted.
- All keys used for comparison must be distinct.
- Each key appears at most twice (one for comparison and as a leaf).


2. Node Tree

- Each node appears only once, together with its object.
- It is preferred for most textbooks because such books don't make the distinction between the keys and the objects

Two Models of Search Tree


1. Leaf Tree

- A tree of height h contains at most 2^h objects (leaves).



2. Node Tree

- A tree of height h contains at most $2^{h+1} - 1$ objects (all nodes).



Two Models of Search Tree

1. Leaf Tree

- It is the preferred in these presentations.
- This is because it requires fewer number of comparisons, and its structure is more regular.
- It will be used for most data structures.

2. Node Tree

- It is used for splay trees.

2. LEAF TREES

Leaf Trees

```
typedef struct tr_n_t {key_t      key;  
                      struct tr_n_t  *left;  
                      struct tr_n_t  *right;  
/* possibly additional information */  
                      } tree_node_t;
```

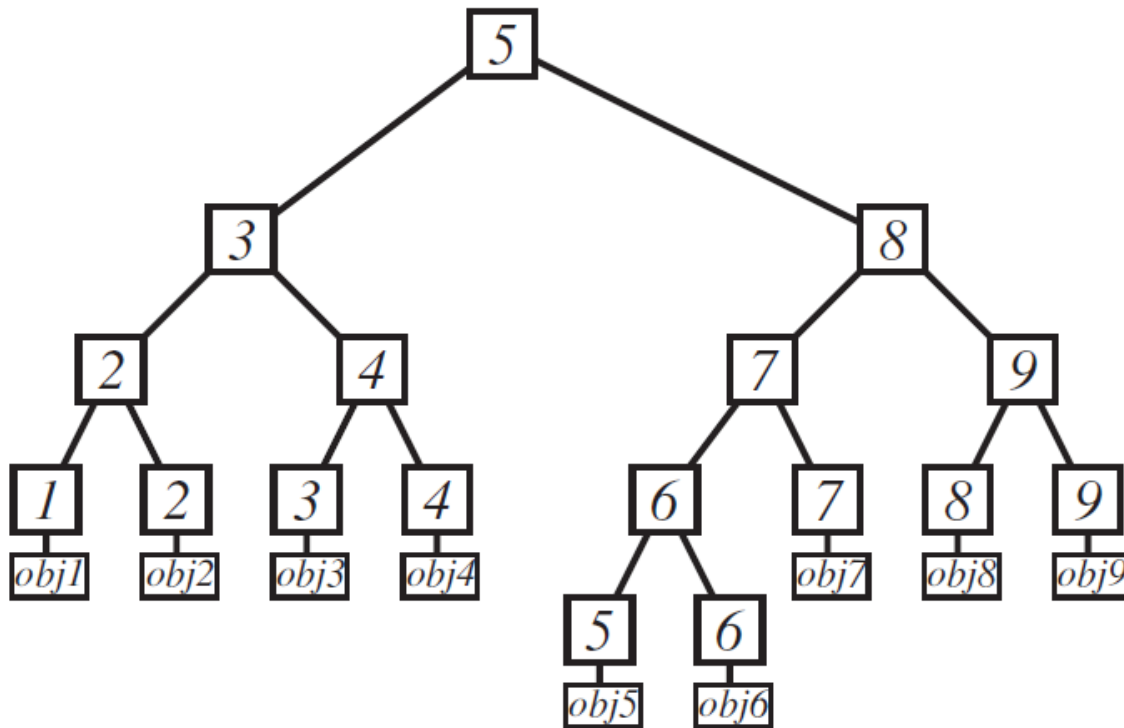
{ A node `*n` is a leaf if `n->right = NULL`. Then `n->left` points to the object stored in that leaf and `n->key` contains the object's key.

{ If `root->left = NULL`, then the tree is empty.

{ If `root->left \neq NULL` and `root->right = NULL`, then `root` is a leaf and the tree contains only one object.

{ If `root->left \neq NULL` and `root->right \neq NULL`, then `root->right` and `root->left` point to the roots of the right and left subtrees. For each node `*left_node` in the left subtree, we have `left_node->key < root->key`, and for each node `*right_node` in the right subtree, we have `right_node->key \geq root->key`.

Leaf Trees




```
tree_node_t *create_tree(void)
{
    tree_node_t *tmp_node;
    tmp_node = get_node();
    tmp_node->left = NULL;
    return( tmp_node );
}
```

Intervals in Leaf Trees

- We can associate each tree node with an interval:


the interval of all possible key values that can be reached through this node.

- The interval of `root` is $]-\infty, \infty[$.
- If $*n$ is an interior node associated with $[a, b[$, then
 - $n \rightarrow \text{key} \in [a, b[$
 - $n \rightarrow \text{left} \in [a, n \rightarrow \text{key}[$
 - $n \rightarrow \text{right} \in [n \rightarrow \text{key}, b[$




Rotations in Leaf Trees

- The same set of (key, object) pairs can be organized in many distinct correct search trees.
- There are two operations that transform a correct search tree in a different correct search tree for the same set.
- They are the left and right rotations.



Rotations in Leaf Trees


```
void left_rotation(tree_node_t *n)
{
    tree_node_t *tmp_node;
    key_t        tmp_key;
    tmp_node = n->left;
    tmp_key  = n->key;
    n->left  = n->right;
    n->key   = n->right->key;
    n->right = n->left->right;
    n->left->right = n->left->left;
    n->left->left  = tmp_node;
    n->left->key   = tmp_key;
}
```



- We move the content of the nodes around, but the node **n* needs to be the root of the subtree because there are pointers from higher levels in the tree that point to **n*.
- This is because there are no pointers to the parent.

Rotations in Leaf Trees

```
void right_rotation(tree_node_t *n)
{
    tree_node_t *tmp_node;
    key_t        tmp_key;
    tmp_node = n->right;
    tmp_key  = n->key;
    n->right = n->left;
    n->key   = n->left->key;
    n->left  = n->right->left;
    n->right->left = n->right->right;
    n->right->right = tmp_node;
    n->right->key   = tmp_key;
}
```



- Any correct search tree for some set of (key, object) pairs can be transformed into each other by sequence of rotations.

Height of a Leaf Tree

- The central property which makes some search trees good and others bad is the height.
- The **height** of a search tree is the maximum length of a path, in number of edges, from the root to a leaf.
- The **depth** of a node is the distance, in number of edges, from the root to such node.

Height of a Leaf Tree

- The maximum number of leaves of a search tree of height h is 2^h .
- The minimum number of leaves is $h + 1$ because
 - a tree of height h must have at least one interior node at each depth $0, \dots, h - 1$.
 - a tree with h interior nodes has $h + 1$ leaves.
- **Theorem:** A leaf tree for n objects has height $\lceil \lg n \rceil \leq h \leq n - 1$.
- **Theorem:** A leaf tree for n objects has average depth of at least $\lg n$ and at most $\frac{(n-1)(n+2)}{2n} \approx \frac{1}{2}n$. Why?

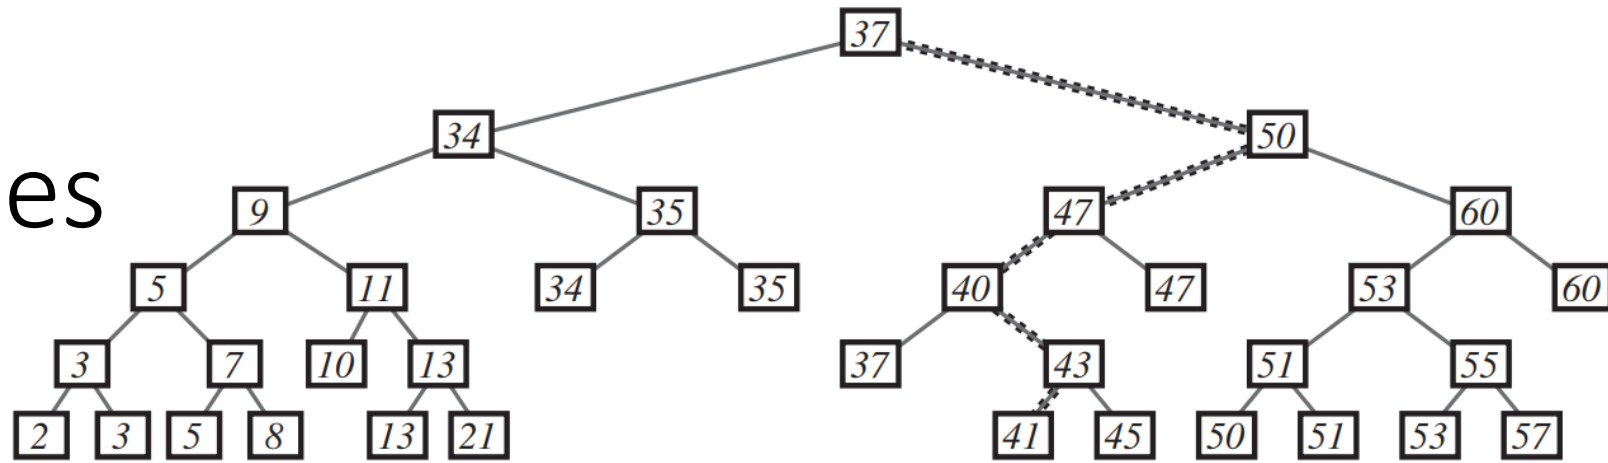
3. OPERATIONS ON LEAF TREES

Operations on Leaf Trees

- { `find(tree, query_key)`: Returns the object associated with `query_key`, if there is one;
- { `insert(tree, key, object)`: Inserts the (key, object) pair in the tree; and
- { `delete(tree, key)`: Deletes the object associated with `key` from the tree.

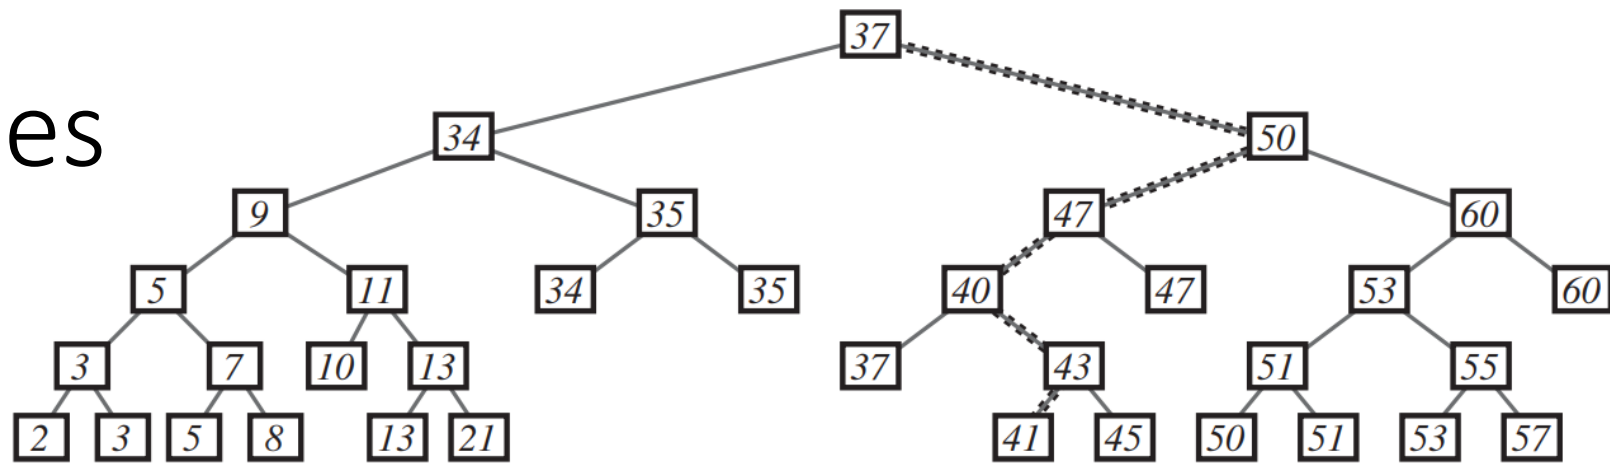
Find on Leaf Trees

```
object_t *find(tree_node_t *tree,
               key_t query_key)
{
    tree_node_t *tmp_node;
    if( tree->left == NULL )
        return(NULL);
    else
    {
        tmp_node = tree;
        while( tmp_node->right != NULL )
        {
            if( query_key < tmp_node->key )
                tmp_node = tmp_node->left;
            else
                tmp_node = tmp_node->right;
        }
        if( tmp_node->key == query_key )
            return( (object_t *) tmp_node->left );
        else
            return( NULL );
    }
}
```



- First, we check if the tree is empty.
- Then, one just follows the associated interval structure to the leaf.
- If the object is there, return it; otherwise return NULL.

Find on Leaf Trees (Recursive)



```
object_t *find(tree_node_t *tree,
               key_t query_key)
{
    if( tree->left == NULL ||
        (tree->right == NULL &&
         tree->key != query_key) )
        return(NULL);
    else if (tree->right == NULL &&
             tree->key == query_key )
        return( (object_t *) tree->left );
    else
    {
        if( query_key < tree->key )
            return( find(tree->left, query_key) );
        else
            return( find(tree->right, query_key) );
    }
}
```

- First, we check if the tree is empty or a leaf that does not contain the `query_key`.
- If it indeed contains the `query_key`, return the object.
- Otherwise, one just follows the associated interval structure to the leaf.

Insert on Leaf Trees

```
int insert(tree_node_t *tree, key_t new_key,
          object_t *new_object)
{
    tree_node_t *tmp_node;
    if( tree->left == NULL )
    {
        tree->left = (tree_node_t *) new_object;
        tree->key   = new_key;
        tree->right  = NULL;
    }
    else
    {
        tmp_node = tree;
        while( tmp_node->right != NULL )
        {
            if( new_key < tmp_node->key )
                tmp_node = tmp_node->left;
            else
                tmp_node = tmp_node->right;
        }
        /* found the candidate leaf. Test whether
           key distinct */
        if( tmp_node->key == new_key )
            return( -1 );
    }
}
```

- If the tree is empty, the new node is added at the root.
- Then, it is similar to find.
- Let us call `tmp_node` as the leaf node where the new key must be inserted.
- It has to create a new interior node and a new leaf node.
- For the moment, assume all the keys are unique.
- Otherwise, it is an error.

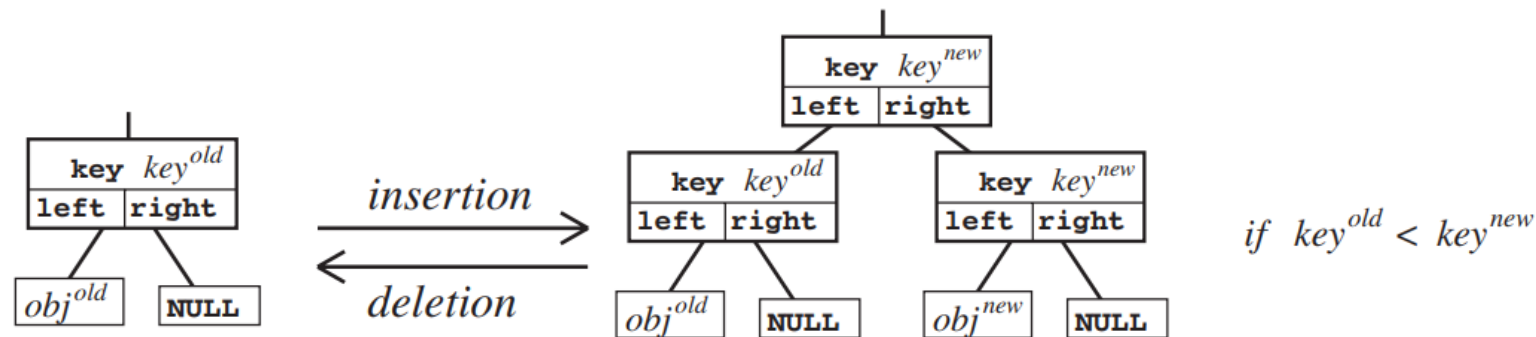
Insert on Leaf Trees

```

/* key is distinct, now perform the insert */
{ tree_node_t *old_leaf, *new_leaf;
  old_leaf = get_node();
  old_leaf->left = tmp_node->left;
  old_leaf->key = tmp_node->key;
  old_leaf->right = NULL;
  new_leaf = get_node();
  new_leaf->left = (tree_node_t *)
    new_object;
  new_leaf->key = new_key;
  new_leaf->right = NULL;
  if( tmp_node->key < new_key )
  { tmp_node->left = old_leaf;
    tmp_node->right = new_leaf;
    tmp_node->key = new_key;
  }
  else
  { tmp_node->left = new_leaf;
    tmp_node->right = old_leaf;
  }
}
return( 0 );

```

- We must create leaf nodes for:
 - former value of tmp_node (old_leaf)
 - new key inserted (new_leaf)
- The node tmp_node, which was a leaf, is now an interior node that points to old_leaf and new_leaf.
- Its key is the largest key of the children.
- The left and right pointers depend on whether $\text{tmp_node} \rightarrow \text{key} < \text{new_leaf} \rightarrow \text{key}$.



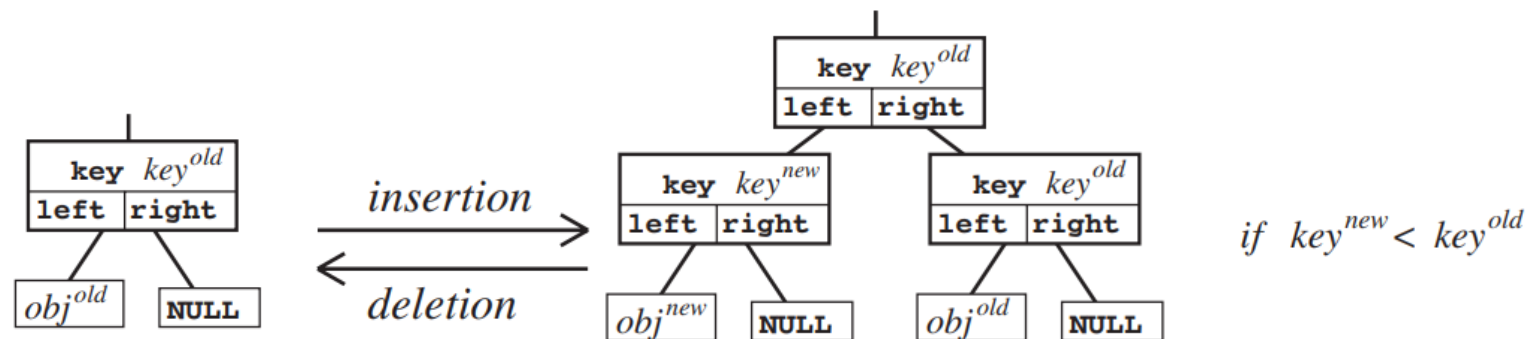
Insert on Leaf Trees

```

/* key is distinct, now perform the insert */
{ tree_node_t *old_leaf, *new_leaf;
  old_leaf = get_node();
  old_leaf->left = tmp_node->left;
  old_leaf->key = tmp_node->key;
  old_leaf->right = NULL;
  new_leaf = get_node();
  new_leaf->left = (tree_node_t *)
  new_object;
  new_leaf->key = new_key;
  new_leaf->right = NULL;
  if( tmp_node->key < new_key )
  { tmp_node->left = old_leaf;
    tmp_node->right = new_leaf;
    tmp_node->key = new_key;
  }
  else
  { tmp_node->left = new_leaf;
    tmp_node->right = old_leaf;
  }
}
return( 0 );

```

- We must create leaf nodes for:
 - former value of tmp_node (old_leaf)
 - new key inserted (new_leaf)
- The node tmp_node, which was a leaf, is now an interior node that points to old_leaf and new_leaf.
- Its key is the largest key of the children.
- The left and right pointers depend on whether $\text{tmp_node} \rightarrow \text{key} < \text{new_leaf} \rightarrow \text{key}$.



Delete on Leaf Trees

```
object_t *delete(tree_node_t *tree,
                  key_t delete_key)
{
    tree_node_t *tmp_node, *upper_node,
    *other_node;
    object_t *deleted_object;
    if( tree->left == NULL )
        return( NULL );
    else if( tree->right == NULL )
    {
        if( tree->key == delete_key )
        {
            deleted_object =
                (object_t *) tree->left;
            tree->left = NULL;
            return( deleted_object );
        }
        else
            return( NULL );
    }
    else
    {
        tmp_node = tree;
```

- If the tree is empty or the key is not in the tree, we return NULL.
- If the root is a leaf and it contains the `delete_key`, we just return it and eliminate the object.
- Otherwise, we must find the corresponding leaf like in `find`.

Delete on Leaf Trees

```
else
{
    tmp_node = tree;
    while( tmp_node->right != NULL )
    {
        upper_node = tmp_node;
        if( delete_key < tmp_node->key )
        {
            tmp_node = upper_node->left;
            other_node = upper_node->right;
        }
        else
        {
            tmp_node = upper_node->right;
            other_node = upper_node->left;
        }
    }
}
```

- Otherwise, we must find the corresponding leaf like in `find`.
- When we are deleting a leaf, we need to delete an interior node above it.
- So, we must keep track of the current node (`tmp_node`) and its upper neighbor (`upper`).
- The sibling of `tmp_node` is `other_node`.
- At the end of the loop, `tmp_node` points to a leaf where `delete_key` may be in.
- But `other_node` may point to a leaf or an interior node.

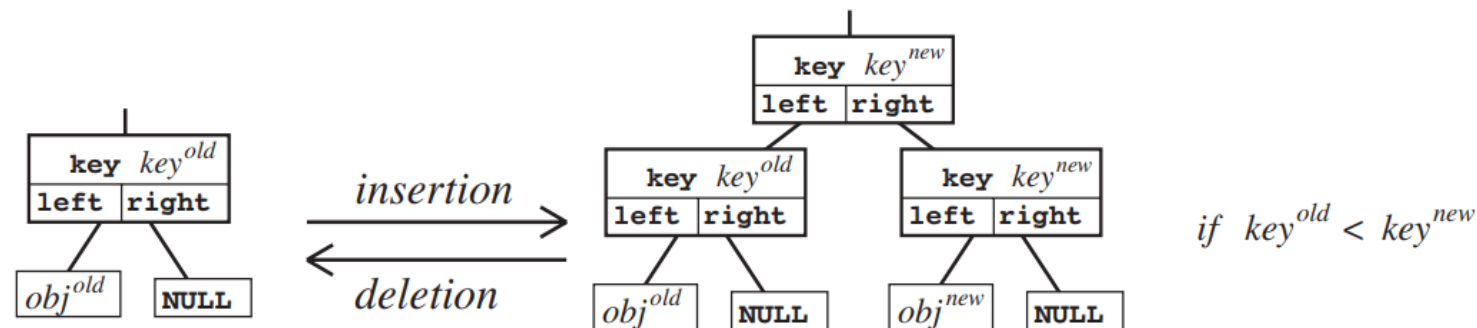
Delete on Leaf Trees

```

if( tmp_node->key != delete_key )
    return( NULL );
else
{
    upper_node->key    = other_node->key;
    upper_node->left   = other_node->left;
    upper_node->right  = other_node->right;
    deleted_object = (object_t *)
        tmp_node->left;
    return_node( tmp_node );
    return_node( other_node );
    return( deleted_object );
}
}

```

- We must store and return the deleted object.
- Also, we must delete a leaf (tmp_node) and an interior node (other_node).
- But before, we copy the key and children of other_node in upper_node.



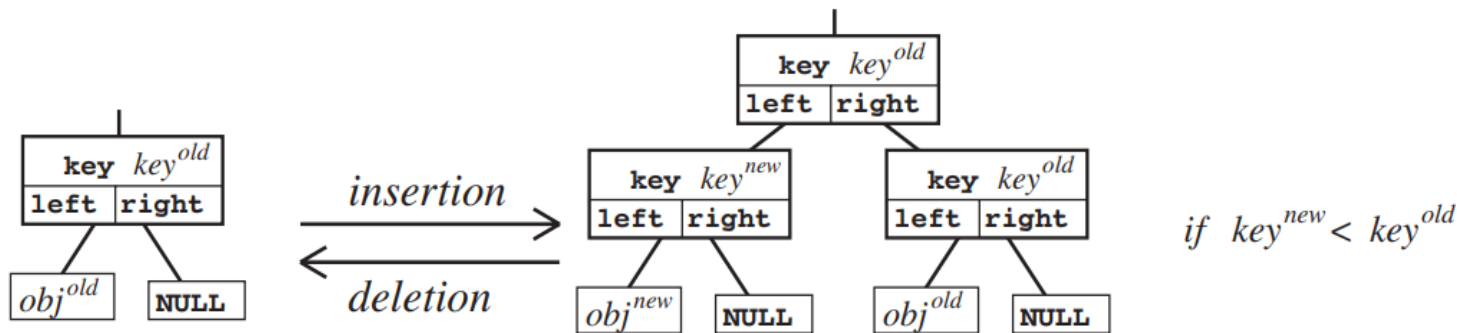
Delete on Leaf Trees

```

if( tmp_node->key != delete_key )
    return( NULL );
else
{
    upper_node->key    = other_node->key;
    upper_node->left   = other_node->left;
    upper_node->right  = other_node->right;
    deleted_object = (object_t *)
        tmp_node->left;
    return_node( tmp_node );
    return_node( other_node );
    return( deleted_object );
}
}

```

- We must store and return the deleted object.
- Also, we must delete a leaf (tmp_node) and an interior node (other_node).
- But before, we copy the key and children of other_node in upper_node.



4. RETURNING FROM LEAF TO ROOT

Returning from Leaf to Root

- We normally start from the root and follow pointers toward a leaf.
- Sometimes, it is necessary to do the opposite. For example, for rebalancing.
- We can do this with any of the following methods:
 1. A stack
 2. Back pointers
 3. Back pointer with lazy update
 4. Reversing the path

1. A Stack

- If we push pointers to all traversed nodes on a stack during descent to a leaf, then we can take the nodes from the stack in the correct reversed order afterwards.
- This is the most economic solution: it does not put additional information into the tree structure.
- The maximum size of the stack needed is the height of the tree.
- This solution is implicitly used in any recursive implementation of the search trees.

2. Back Pointers

- If each node contains also a pointer to its parent, then we have a path up from any node back to the root.
- It requires additional memory. However, now it is not a problem.
- However, this pointer has to be corrected in each operation. This requires time and may generate errors.

3. Back Pointer with Lazy Update

- Each node has a pointer to its parent.
- However, it is only entered during descent in the tree.
- Then, we have a correct path from the leaf we reached to the root.
- We don't need to correct back pointers during all operations of the tree.
- But then, a back pointer is only assumed to be correct for the nodes on the path from the leaf we just reached.

4. Reversing the Path

- We can keep back pointers for the path even without an extra entry for a back pointer.
- We just reverse the forward pointers as we go down the tree.
- For instance, if we go left, the left pointer is used as a back pointer.
- When we go up again, the correct forward pointers must be restored.
- This method does not use extra space.
- But it causes many additional problems because the leaf tree structure is temporarily destroyed.
- It is not recommended to use it.

5. DEALING WITH NONUNIQUE KEYS

Dealing with Nonunique Keys

- In practical applications, it is not uncommon that there are several objects with the same key.
- In database applications, there are queries to list all objects with a given attribute value.
- Thus, we must adapt the `find`, `insert` and `delete` operations.

Dealing with Nonunique Keys

- { `find` returns all objects whose key is the given query key in output-sensitive time $O(h + k)$, where h is the height of the tree and k is the number of elements that `find` returns.
- { `insert` always inserts correctly in time $O(h)$, where h is the height of the tree.
- { `delete` deletes all items of that key in time $O(h)$, where h is the height of the tree.

Dealing with Nonunique Keys

The obvious way to realize this behavior is to keep all elements of the same key in a linked list below the corresponding leaf of the leaf tree.

- `find` just returns all the elements of that list.
- `insert` always inserts at the beginning of that list.
- `delete` requires additional information.
 - We need an additional node, between the leaf and the linked list, which contains pointers to the beginning and the end of the list.
 - Then, we can transfer the entire list with $O(1)$ operations to the free list of our dynamic memory allocation structure.

6. QUERIES FOR THE KEYS IN AN INTERVAL

Queries for the Keys in an Interval

- If the keys are subject to small errors, we might not know the key exactly, so we may want to look for an interval of keys instead.
- We give an interval $[a,b[$ and want to find all keys that are contained in this interval.
- This can be adapted in several ways:
 1. Organize the leaves in a doubly linked list.
 2. Change the query and not the data structure.

1. Leaves in a Doubly Linked List

- We can move in $O(1)$ time from a leaf to the next larger and the next smaller leaf.
- This requires a change in the insertion and deletion functions to maintain the list.
- But it is an easy change that takes only $O(1)$ additional time. How?
- The query method is also almost the same. It takes $O(k)$ additional time if it lists a total of k keys in the interval.

2. Change in the Query and not in the Tree

- We go down the tree with the query interval instead of the query key:
 - We go left if $[a, b[< \text{node} \rightarrow \text{key}$.
 - We go right if $\text{node} \rightarrow \text{key} \leq [a, b[$.
 - We go both left and right if $a < \text{node} \rightarrow \text{key} \leq b$.
- We store all the branches that we still need to explore on a stack. The nodes we visit this way are:
 - Nodes on the search path for a .
 - Nodes on the search path for b .
 - Nodes on the tree between these paths.

2. Change in the Query and not in the Tree

- If there are i interior nodes between these paths, there must be at least $i + 1$ leaves between these paths.
- So if this method lists k leaves, the total number of nodes visited is at most twice the number of nodes visited in a normal find operation plus $O(k)$. Why?
- This method is slightly slower than the first one, but it requires no change in the `insert` and `delete` operations.

2. Implementation

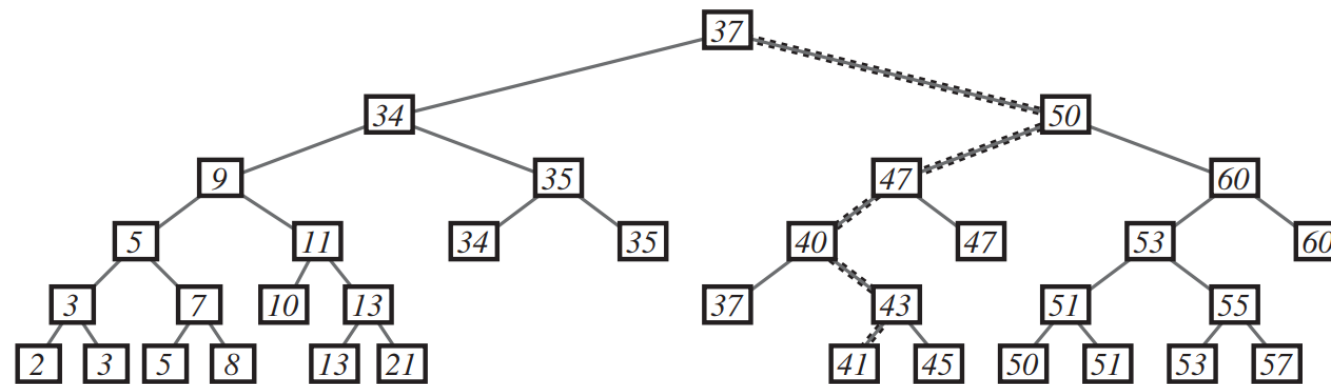
- The output of the operation is potentially long.
- Thus, we return many objects instead of a single result.
- For that, we create a linked list of the (key, object) pairs found in the query interval, which is linked by the right pointers.
- In particular, each node of such list is a tree node `node`, where the key is stored in `node->key` and the object is stored in `node->right`.

2. Implementation

```

tree_node_t *interval_find(tree_node_t *tree,
                           key_t a, key_t b)
{
    tree_node_t *tr_node;
    tree_node_t *result_list, *tmp;
    result_list = NULL;
    create_stack();
    push(tree);
    while( !stack_empty() )
    {
        tr_node = pop();
        if( tr_node->right == NULL )
        {
            /* reached leaf, now test */
            if( a <= tr_node->key &&
                tr_node->key < b )
            {
                tmp = get_node();
                /* leaf key in interval */
                tmp->key = tr_node->key; /*
                copy to output list */
                tmp->left = tr_node->left;
                tmp->right = result_list;
                result_list = tmp;
            }
        }
        /* not leaf, might have to follow down */
    }
}

```




2. Implementation

```

} /* not leaf, might have to follow down */
else if ( b <= tr_node->key )
/* entire interval left */
    push( tr_node->left );
else if ( tr_node->key <= a )
/* entire interval right */
    push( tr_node->right );
else /* node key in interval,
        follow left and right */
{
    push( tr_node->left );
    push( tr_node->right );
}
}
remove_stack();
return( result_list );

```



7. OPTIMAL SEARCH TREES


Optimal Search Trees

- We consider search trees as a static structure: no inserts and no deletes.
- Then, there is no problem of rebalancing the tree.
- We should build it as good as possible: minimum height.
- Because a tree of height h has at most 2^h , an optimal search tree for n items has height $\lceil \lg n \rceil$.
- There are two ways to construct it:
 1. Bottom up
 2. Top Down

```
typedef struct tr_n_t {key_t      key;
                      struct tr_n_t *left;
                      struct tr_n_t *right;
                      /* possibly additional information */
                      } tree_node_t;
```

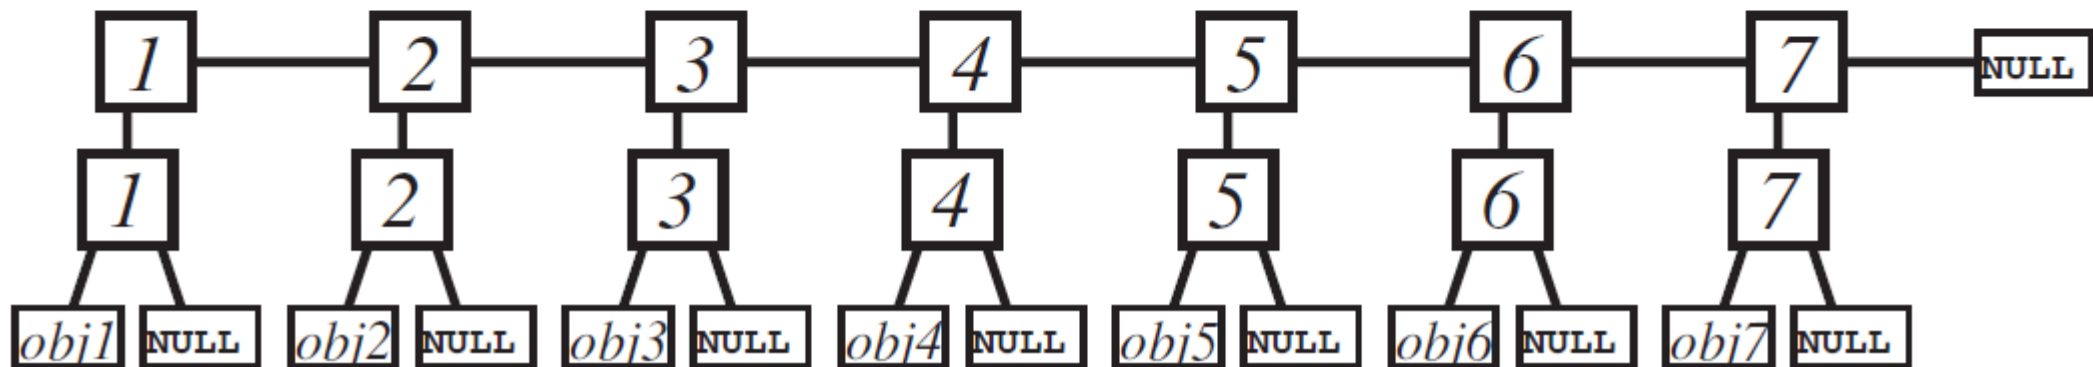

Bottom Up Construction

- The input is a sorted list of the (key, object pairs).
- In particular, it is a list of `tree_node_t` nodes where each node contains:
 - `key`
 - `left` points to the object
 - `right` points to the next element (or NULL at the end of the list).



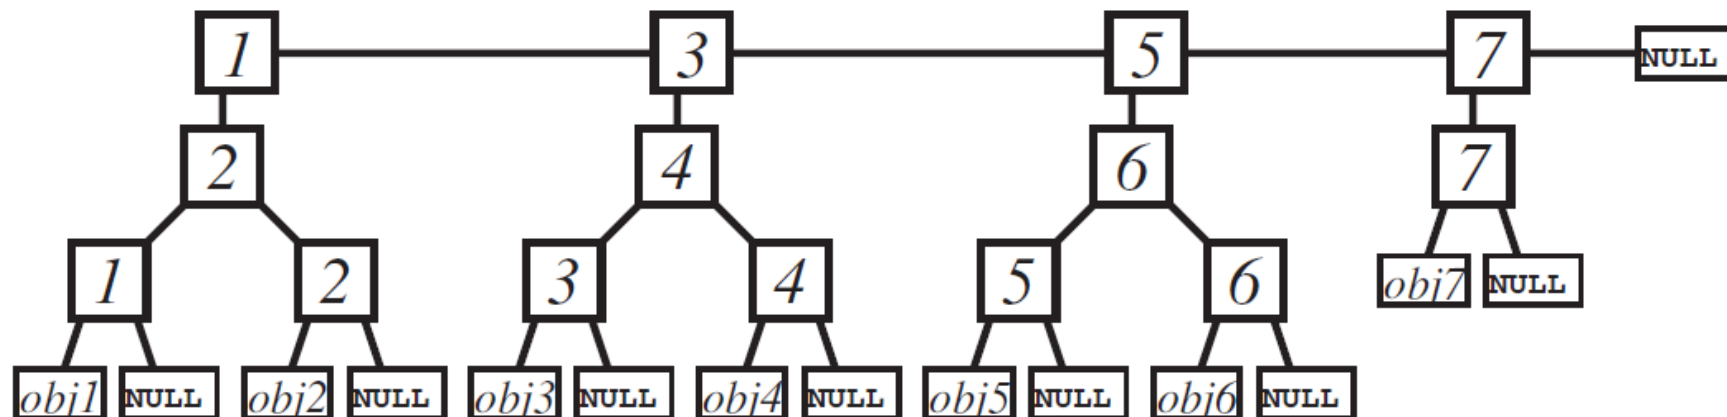
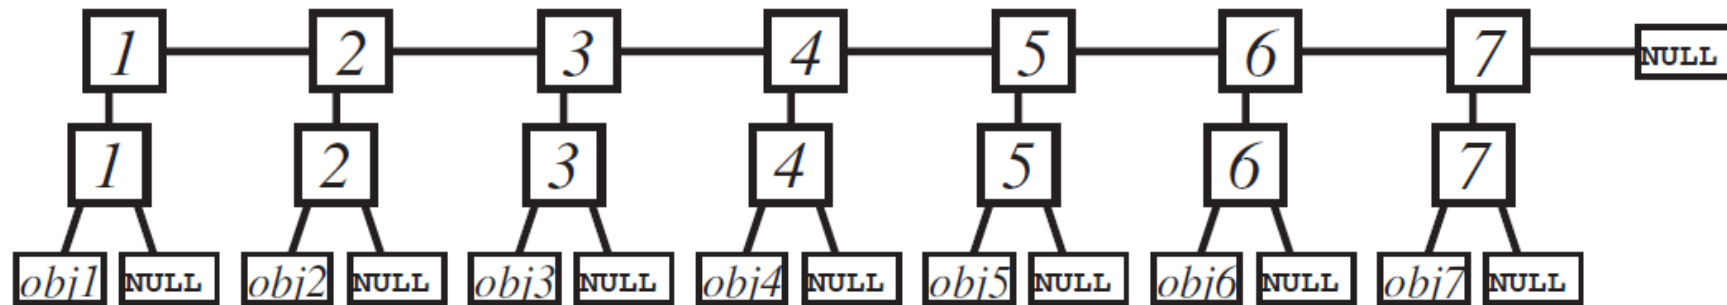
Bottom Up Construction

- Create a list as a list of one-element trees (leaves).



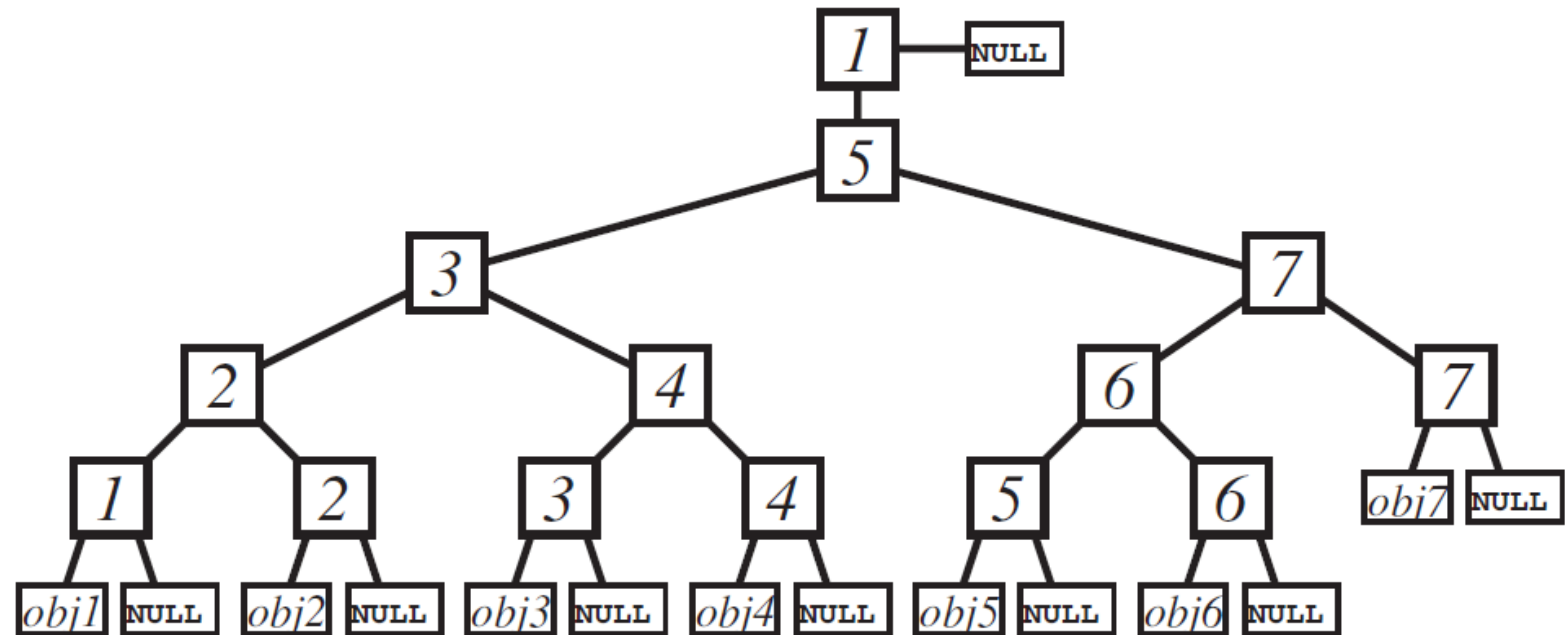
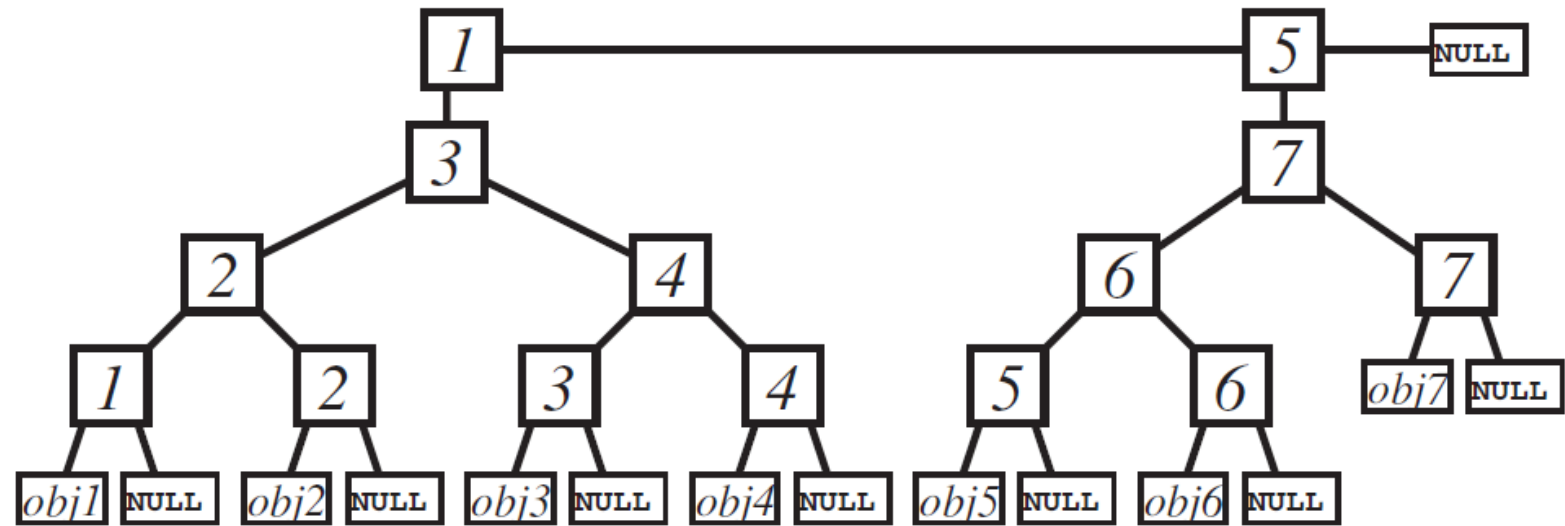
Bottom Up Construction

- Go through the list , joining two consecutive trees.
- The nodes of the previous list are attached as leaves.





Bottom Up

- Repeat the process until there is only one tree left.



Bottom Up



- If the list is empty, create an empty tree.
- Otherwise, create a list as a list of one-element trees (leaves).



```
tree_node_t *make_tree(tree_node_t *list)
{
    tree_node_t *end, *root;
    if( list == NULL )
    {
        root = get_node(); /* create empty tree */
        root->left = root->right = NULL;
        return( root );
    }
    else if( list->right == NULL )
        return( list ); /* one-leaf tree */
    else /* nontrivial work required: at least
          two nodes */
    {
        root = end = get_node();
        /* convert input list into leaves below
           new list */
        end->left = list;
        end->key = list->key;
        list = list->right;
        end->left->right = NULL;
        while( list != NULL )
        {
            end->right = get_node();
            end = end->right;
            end->left = list;
            end->key = list->key;
            list = list->right;
            end->left->right = NULL;
        }
        end->right = NULL;
    }
}
```

Bottom Up

- We create the upper part of the new list with pointer `end`, i.e. the nodes that point to the leaves.



```
tree_node_t *make_tree(tree_node_t *list)
{
    tree_node_t *end, *root;
    if( list == NULL )
    {
        root = get_node(); /* create empty tree */
        root->left = root->right = NULL;
        return( root );
    }
    else if( list->right == NULL )
        return( list ); /* one-leaf tree */
    else /* nontrivial work required: at least
        two nodes */
    {
        root = end = get_node();
        /* convert input list into leaves below
        new list */
        end->left = list;
        end->key = list->key;
        list = list->right;
        end->left->right = NULL;
        while( list != NULL )
        {
            end->right = get_node();
            end = end->right;
            end->left = list;
            end->key = list->key;
            list = list->right;
            end->left->right = NULL;
        }
        end->right = NULL;
    }
}
```



```

{ tree_node_t *old_list, *new_list, *tmp1,
  *tmp2;

old_list = root;
while( old_list->right != NULL )
{ /* join first two trees from
   old_list */
  tmp1 = old_list;
  tmp2 = old_list->right;
  old_list = old_list->right->right;
  tmp2->right = tmp2->left;
  tmp2->left = tmp1->left;
  tmp1->left = tmp2;
  tmp1->right = NULL;
  new_list = end = tmp1;
  /* new_list started */
  while( old_list != NULL )
  /* not at end */
  { if( old_list->right == NULL )
    /* last tree */
    { end->right = old_list;
      old_list = NULL;
    }
  }
}

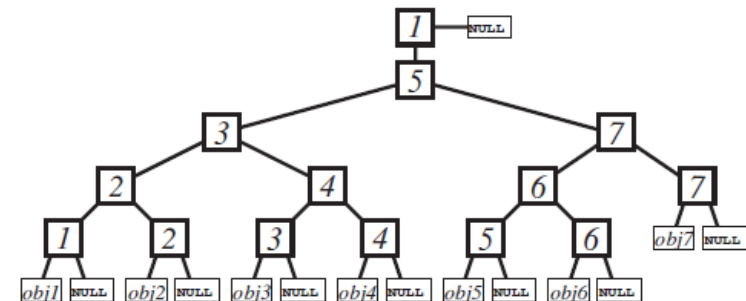



```

```

else /* join next two trees of
      old_list */
{ tmp1 = old_list;
  tmp2 = old_list->right;
  old_list =
    old_list->right->right;
  tmp2->right = tmp2->left;
  tmp2->left = tmp1->left;
  tmp1->left = tmp2;
  tmp1->right = NULL;
  end->right = tmp1;
  end = end->right;
}
} /* finished one pass through
   old_list */
old_list = new_list;
} /* end joining pairs of trees
   together */
root = old_list->left;
return_node( old_list );
}
return( root );
}

```

- The inner while goes through the list joining two consecutive trees.
- The outer while repeats the process until there is only one tree left.

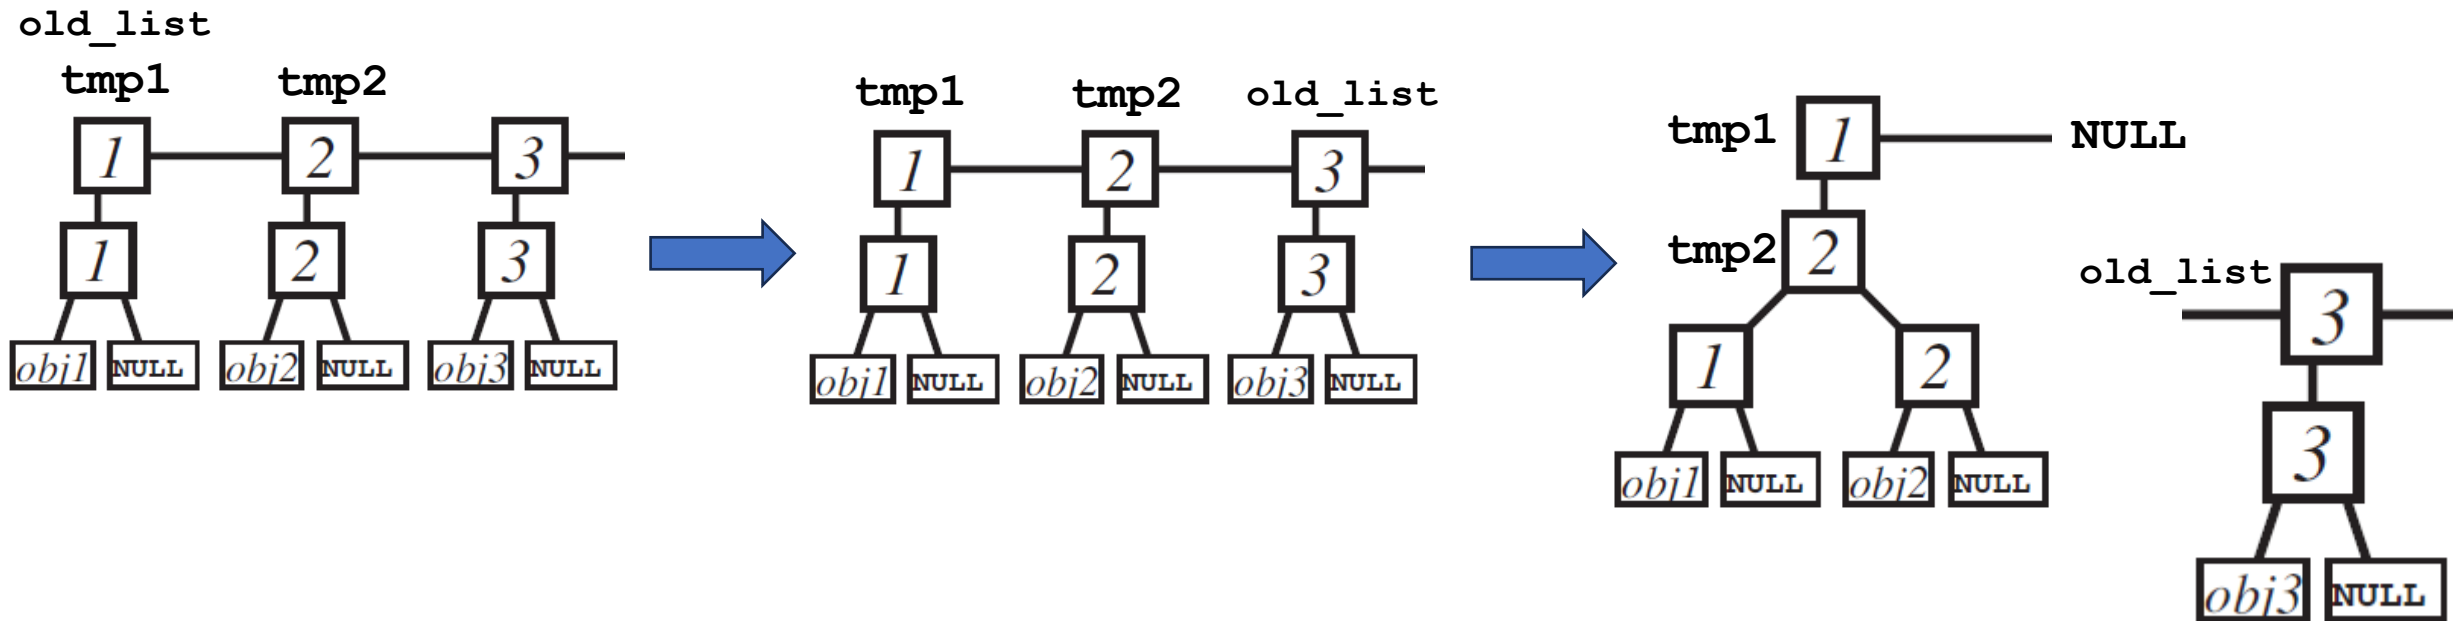


- In each iteration of the outer while, `old_list` is used to construct a `new_list` that merges the consecutive pairs of trees of `old_list`.
- During the process, `new_list` points to what has been merged until `end`; `old_list` points to what is left to be merged.

Merging two consecutive trees

- In each iteration of the outer while, `old_list` is used to construct a `new_list` that merges the consecutive pairs of trees of `old_list`.
- During the process, `new_list` points to what has been merged until end; `old_list` points to what is left to be merged.
- The objects that are to be merged are `tmp1` and `tmp2`.

```
tmp1 = old_list;
tmp2 = old_list->right;
old_list = old_list->right->right;
tmp2->right = tmp2->left;
tmp2->left = tmp1->left;
tmp1->left = tmp2;
tmp1->right = NULL;
```




Merging two consecutive trees

This operation is performed:

- at the beginning of each iteration of the outer loop.
 - We keep the beginning of the list in `new_list`.
- when there are still trees to merge in the inner loop.
 - We append the merged nodes into `new_list` at the end and update `end`.

```
tmp1 = old_list;  
tmp2 = old_list->right;  
old_list = old_list->right->right;  
tmp2->right = tmp2->left;  
tmp2->left = tmp1->left;  
tmp1->left = tmp2;  
tmp1->right = NULL;
```



```

{ tree_node_t *old_list, *new_list, *tmp1,
  *tmp2;

old_list = root;
while( old_list->right != NULL )
{ /* join first two trees from
   old_list */
  tmp1 = old_list;
  tmp2 = old_list->right;
  old_list = old_list->right->right;
  tmp2->right = tmp2->left;
  tmp2->left = tmp1->left;
  tmp1->left = tmp2;
  tmp1->right = NULL;
  new_list = end = tmp1;
  /* new_list started */
  while( old_list != NULL )
  /* not at end */
  { if( old_list->right == NULL )
    /* last tree */
    { end->right = old_list;
      old_list = NULL;
    }
  }
}

```





```

else /* join next two trees of
      old_list */
{ tmp1 = old_list;
  tmp2 = old_list->right;
  old_list =
    old_list->right->right;
  tmp2->right = tmp2->left;
  tmp2->left = tmp1->left;
  tmp1->left = tmp2;
  tmp1->right = NULL;
  end->right = tmp1;
  end = end->right;
}
} /* finished one pass through
   old_list */
old_list = new_list;
} /* end joining pairs of trees
   together */
root = old_list->left;
return_node( old_list );
}
return( root );
}

```

The merge operation is performed:

- at the beginning of each iteration of the outer loop.
 - We keep the beginning of the list in `new_list`.



- when there are still trees to merge in the inner loop.
 - We append the merged nodes into `new_list` at the end and update `end`.

```
{ tree_node_t *old_list, *new_list, *tmp1,
    *tmp2;
```

```
old_list = root;
while( old_list->right != NULL )
{ /* join first two trees from
    old_list */
```

```
tmp1 = old_list;
tmp2 = old_list->right;
old_list = old_list->right->right;
tmp2->right = tmp2->left;
tmp2->left = tmp1->left;
tmp1->left = tmp2;
tmp1->right = NULL;
new_list = end = tmp1;
/* new_list started */
while( old_list != NULL )
/* not at end */
{ if( old_list->right == NULL )
    /* last tree */
    { end->right = old_list;
      old_list = NULL;
    }
}
```

```
else /* join next two trees of
    old_list */
{ tmp1 = old_list;
  tmp2 = old_list->right;
  old_list =
      old_list->right->right;
  tmp2->right = tmp2->left;
  tmp2->left = tmp1->left;
  tmp1->left = tmp2;
  tmp1->right = NULL;
  end->right = tmp1;
  end = end->right;
}
} /* finished one pass through
    old_list */
old_list = new_list;
} /* end joining pairs of trees
    together */
root = old_list->left;
return_node( old_list );
}
return( root );
}
```

We have a special case when there is an element that does not have another node to be merged with.

- We just add it as the final node of `new_list`.

- At the end of an iteration of the outer loop, we just set the new list as the old one.
- When the process is finished, we return the resulting tree and free the list that points to it.

Complexity Analysis of the Bottom Up Construction

- This method constructs a search tree of optimal height from an ordered list in time $O(n)$.
- **First Part:** Duplicating the list and converting it in a list of leaves takes $O(n)$.
- **Second Part:** In each iteration of the innermost while, one of the n interior nodes created in the first part is removed from the current list and put into a finished subtree. Thus, it is executed only $O(n)$ times.

Top Down Construction



- Recursive description
 - Divide the data set in the middle.
 - Create optimal trees for the lower and the upper halves.
 - Join them together.
- This division is balanced: the number of items in the left and in the right subtree differs by at most one.
- It results in a tree of optimal height.
- But if we implement it this way, the complexity is $\Theta(n \lg n)$ due to the $O(n)$ overhead in each recursion to find the middle of the list.
- But we can implement it in $O(n)$ using a stack. The best option is an array-based stack.

Top Down Construction

- We first construct the tree “in the abstract” without filling any key values or pointers to objects.
- We don’t need to find the middle of the list; we just need to keep track of the number of elements that should go into the left and right subtrees.
- We can build this abstract tree of the required shape easily using a stack.
- We initially put the root on the stack labeled with the required tree size.
- Then, we continue until the stack is empty, to
 - Pop a node from the stack.
 - If its size is 1 (it’s a leaf), it should be filled with the next key and object from the list.
 - Otherwise,
 - Attach it to two newly created nodes labeled with half the size.
 - Put the new nodes on the stack.

Top Down Construction

- The problem is to fill in the keys of the interior nodes, which become available only when a leaf is reached.
- For this, each item on the stack needs two pointers:
 - `node1`: The node that still needs to be expanded. The one we are considering.
 - `node2`: The node higher up in the tree, where the smallest key of leaves below that node's right subtree should be inserted as comparison key.



Top Down Construction

- The problem is to fill in the keys of the interior nodes, which become available only when a leaf is reached.
- For this, each item on the stack needs two pointers:
 - `node1`: The node that still needs to be expanded. The one we are considering.
 - `node2`: The node higher up in the tree, where the smallest key of leaves below that node's right subtree should be inserted as comparison key.
- Also, each stack item contains a `number`: the number of leaves that should be created below that node.
- When we pop a node from the stack and create its two children, we first insert the right child first.
- Then, when we reach a leaf, it is the leftmost unfinished leaf of the tree.

Top Down Construction

- The problem is to fill in the keys of the interior nodes, which become available only when a leaf is reached.
- For this, each item on the stack needs two pointers:
 - `node1`: The node that still needs to be expanded. The one we are considering.
 - `node2`: The node higher up in the tree, where the smallest key of leaves below that node's right subtree should be inserted as comparison key.
- The pointer for the missing key value propagates into the left subtree of the current node (where the smallest node comes from).
- The smallest key from the right subtree should become the comparison key of the current node.

Top Down Construction

- The entries for a stack item are:
 - `node1`: The node that still needs to be expanded. The one we are considering.
 - `node2`: Pointer to the node higher up in the tree, where the smallest key of leaves below that node's right subtree should be inserted as comparison key.
 - `number`: the number of leaves that should be created below that node.
- First, we count the number of nodes.
- We create the root and its corresponding stack item. Then, we push it into the stack.

```
tree_node_t *make_tree(tree_node_t *list)
{
    typedef struct { tree_node_t  *node1;
                     tree_node_t  *node2;
                     int           number; }
                     st_item;

    st_item current, left, right;
    tree_node_t *tmp, *root;
    int length = 0;
    for( tmp = list; tmp != NULL;
        tmp = tmp->right )
        length += 1; /* find length of list */
    create_stack(); /* stack of st_item:
                     replace by array */
    root = get_node();
    /* put root node on stack */
    current.node1 = root;
    current.node2 = NULL;
    /* root expands to length leaves */
    current.number = length;
    push( current );
}
```


Top Down Construction

- When we pop a stack item `current` and it does not correspond to a leaf, we create its children `left` and `right` with their corresponding sizes.
- `node1` corresponds to each created child that may need to be expanded.
- `node2` is set to:
 - `current.node2` (for `left`)
 - `current` (for `right`)
- We set the pointers between `current` and `left` and `right`.
- Then, we push `right` and `left` into the stack.

```
while( !stack_empty() )
/* there is still unexpanded node */
{  current = pop();
   if( current.number > 1 )
   /* create (empty) tree nodes */
   {  left.node1 = get_node();
      left.node2 = current.node2;
      left.number = current.number / 2;
      right.node1 = get_node();
      right.node2 = current.node1;
      right.number = current.number -
                      left.number;
      (current.node1)->left  = left.node1;
      (current.node1)->right = right.node1;
      push( right );
      push( left );
   }
   else /* reached a leaf, must be filled
        with list item */
```

Top Down Construction

- node2 is set to:
 - `current.node2` (for left)
 - **`current` (for right)**
- Note that for the right child, node2 points to its parent `current`.
- Notice that `current` is the closest ancestor that is to the left of `right`.
- Thus, if `right` is a leaf, its key should be the comparison value of its parent `current`.




```
while( !stack_empty() )
/* there is still unexpanded node */
{  current = pop();
  if( current.number > 1 )
  /* create (empty) tree nodes */
  { left.node1 = get_node();
    left.node2 = current.node2;
    left.number = current.number / 2;
    right.node1 = get_node();
    right.node2 = current.node1;
    right.number = current.number -
                    left.number;

    (current.node1)->left  = left.node1;
    (current.node1)->right = right.node1;
    push( right );
    push( left );
  }
  else /* reached a leaf, must be filled
        with list item */
```

Top Down Construction

- node2 is set to:
 - `current.node2` (for left)
 - **current** (for right)
- All the internal nodes whose right children are leaves are getting a comparison value.




```
while( !stack_empty() )
/* there is still unexpanded node */
{
    current = pop();
    if( current.number > 1 )
    /* create (empty) tree nodes */
    {
        left.node1 = get_node();
        left.node2 = current.node2;
        left.number = current.number / 2;
        right.node1 = get_node();
        right.node2 = current.node1;
        right.number = current.number -
                        left.number;

        (current.node1)->left  = left.node1;
        (current.node1)->right = right.node1;
        push( right );
        push( left );
    }
    else /* reached a leaf, must be filled
          with list item */

```


Top Down Construction


- node2 is set to:
 - **current.node2** (for left)
 - current (for right)
- Note that for the left child, node2 may point to either NULL or a non immediate ancestor.
- If it points to NULL it is because left is the leftmost node so far.



```
while( !stack_empty() )
/* there is still unexpanded node */
{  current = pop();
  if( current.number > 1 )
  /* create (empty) tree nodes */
  { left.node1 = get_node();
    left.node2 = current.node2;
    left.number = current.number / 2;
    right.node1 = get_node();
    right.node2 = current.node1;
    right.number = current.number -
                    left.number;
    (current.node1)->left  = left.node1;
    (current.node1)->right = right.node1;
    push( right );
    push( left );
  }
  else /* reached a leaf, must be filled
        with list item */
```

Top Down Construction

- `node2` is set to:
 - **`current.node2`** (for `left`)
 - `current` (for `right`)
- Otherwise, it points to the closest ancestor that is to the left of `left`.
- Thus, if `left` is a leaf, its key should be the comparison value of `current.node2`.




```
while( !stack_empty() )
/* there is still unexpanded node */
{
    current = pop();
    if( current.number > 1 )
    /* create (empty) tree nodes */
    {
        left.node1 = get_node();
        left.node2 = current.node2;
        left.number = current.number / 2;
        right.node1 = get_node();
        right.node2 = current.node1;
        right.number = current.number -
                        left.number;
        (current.node1)->left  = left.node1;
        (current.node1)->right = right.node1;
        push( right );
        push( left );
    }
    else /* reached a leaf, must be filled
        with list item */

```

Top Down Construction

- node2 is set to:
 - **current.node2** (for left)
 - current (for right)
- Otherwise, it points to the closest ancestor that is to the left of left.
- Note that node2 was propagated to the left until it reaches a leaf.



```
while( !stack_empty() )
/* there is still unexpanded node */
{
    current = pop();
    if( current.number > 1 )
    /* create (empty) tree nodes */
    {
        left.node1 = get_node();
        left.node2 = current.node2;
        left.number = current.number / 2;
        right.node1 = get_node();
        right.node2 = current.node1;
        right.number = current.number -
                        left.number;

        (current.node1)->left  = left.node1;
        (current.node1)->right = right.node1;
        push( right );
        push( left );
    }
    else /* reached a leaf, must be filled
        with list item */

```

Top Down Construction

- node2 is set to:
 - **current.node2** (for left)
 - current (for right)
- Note that for the left child, node2 may point to a non immediate ancestor, i.e. the closest ancestor that is to the left of left.
- Thus, if left is a leaf, its key should be the comparison value of current.node2.
- All the internal nodes are getting a comparison value because there is always a leaf for which it is the closest ancestor to the left of it.

```
while( !stack_empty() )
/* there is still unexpanded node */
{
    current = pop();
    if( current.number > 1 )
    /* create (empty) tree nodes */
    {
        left.node1 = get_node();
        left.node2 = current.node2;
        left.number = current.number / 2;
        right.node1 = get_node();
        right.node2 = current.node1;
        right.number = current.number -
                        left.number;
        (current.node1)->left = left.node1;
        (current.node1)->right = right.node1;
        push( right );
        push( left );
    }
    else /* reached a leaf, must be filled
          with list item */






```

Top Down Construction







- When we pop a stack item `current` and it corresponds to a leaf, we fill the corresponding information.
- If `current.node2` is not NULL, we fill the comparison value of such node.
- We remove the inserted element from the list and move on to the next element.
- After the loop is completed, the root of the tree is returned.

```
else /* reached a leaf, must be filled
      with list item */
{ (current.node1)->left  = list->left;
  /* fill leaf from list */
  (current.node1)->key   = list->key;
  (current.node1)->right = NULL;
  if( current.node2 != NULL )
    /* insert comparison key in
       interior node */
    (current.node2)->key = list->key;
  tmp = list;
  /* unlink first item from list */
  list = list->right;
  /* content has been copied to */
  return_node(tmp);
  /* leaf, so node is returned */
}
}
return( root );
}
```





Top Down Construction




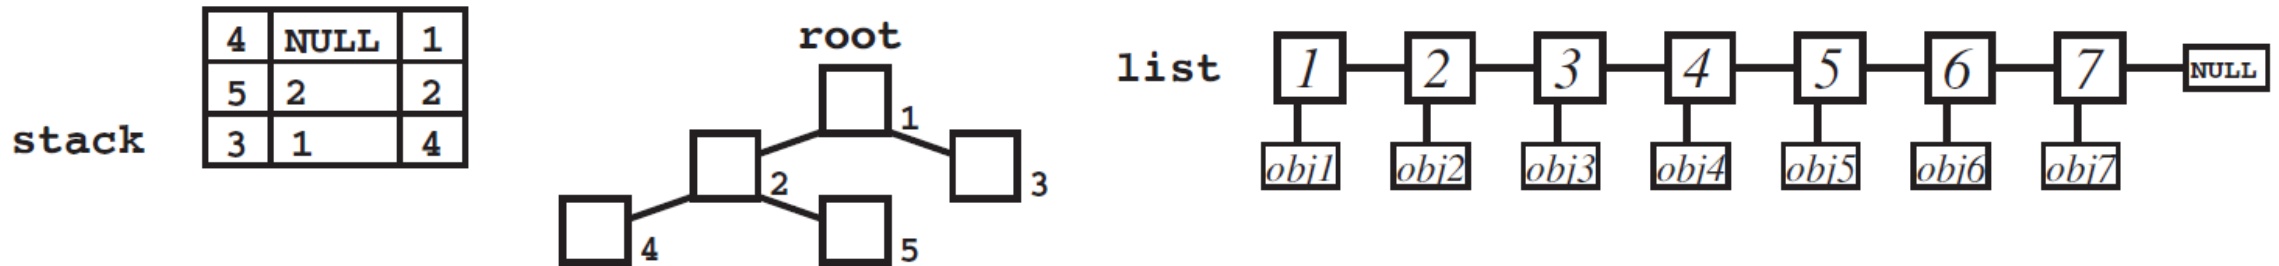

Top Down Construction



Top Down Construction



Top Down Construction





Complexity Analysis of the Top Down Construction

- This method constructs a search tree of optimal height from an ordered list in $O(n)$ time.
- Observe that in each step on the stack we either
 - create two new nodes, and there are only $n - 1$ nodes created in total, or
 - we attach a list item as a leaf, and there are only n list items.
- There are other methods to construct the top-down optimal tree.
- They differ only in the amount of additional space needed.
- In this algorithm, it is the size of the stack: $\lceil \lg n \rceil$.

Analysis of the Top Down Construction

- It is more complicated of the bottom up construction, but it constructs a more balanced tree.



BOTTOM-UP AND TOP-DOWN OPTIMAL TREE WITH 18 LEAVES

8. CONVERTING A TREE INTO A SORTED LIST

Converting a Tree into a Sorted List

- We can use a stack for a trivial depth-first search enumeration of the leaves in decreasing order.
- This is because the right children will be popped first.
- We insert each node in the front of the list.
- Thus, we obtain an increasing order list of elements.
- The size of the stack is the height of the tree.
- It takes $O(n)$ time.

```
tree_node_t *make_list(tree_node_t *tree)
{
    tree_node_t *list, *node;
    if( tree->left == NULL )
    {
        return_node( tree );
        return( NULL );
    }
    else
    {
        create_stack();
        push( tree );
        list = NULL;
        while( !stack_empty() )
        {
            node = pop();
            if( node->right == NULL )
            {
                node->right = list;
                list = node;
            }
            else
            {
                push( node->left );
                push( node->right );
                return_node( node );
            }
        }
        return( list );
    }
}
```


9. REMOVING A TREE

Removing a Tree

- It is important to free all nodes to avoid a memory leak.
- We can do it in $O(n)$ time, i.e. constant time per freed node.
- We can do it using a stack like in the previous algorithm.
- Alternatively, we can perform right rotations in the root until the left-lower neighbor is a leaf.
- Then, it returns the leaf, moves the root down to the right and return the previous root.

```
void remove_tree(tree_node_t *tree)
{
    tree_node_t *current_node, *tmp;
    if( tree->left == NULL )
        return_node( tree );
    else
    {
        current_node = tree;
        while(current_node->right != NULL )
        {
            if( current_node->left->right == NULL )
            {
                return_node( current_node->left );
                tmp = current_node->right;
                return_node( current_node );
                current_node = tmp;
            }
            else
            {
                tmp = current_node->left;
                current_node->left = tmp->right;
                tmp->right = current_node;
                current_node = tmp;
            }
        }
        return_node( current_node );
    }
}
```

Removing a Tree



```
void remove_tree(tree_node_t *tree)
{
    tree_node_t *current_node, *tmp;
    if( tree->left == NULL )
        return_node( tree );
    else
    {
        current_node = tree;
        while(current_node->right != NULL )
        {
            if( current_node->left->right == NULL )
            {
                return_node( current_node->left );
                tmp = current_node->right;
                return_node( current_node );
                current_node = tmp;
            }
            else
            {
                tmp = current_node->left;
                current_node->left = tmp->right;
                tmp->right = current_node;
                current_node = tmp;
            }
        }
        return_node( current_node );
    }
}
```


BIBLIOGRAPHY

- Peter Brass. Advanced Data Structures. Cambridge University Press. 2008.